

November_Rain Presentation



Group Members:

Samuel Colon

Xiaoya Deng

Jonathan Driscoll





Step 1

- Read .json files into Spark RDD

Step 2

- Build a base model – to describe the closing price of a company's stock as a function of relevant word frequency

Step 3

- Improve the base model using sentiment analysis



- In the process of **building a simple model** which describe the closing price of a company's stock as a function of how many times that company has been mentioned on twitter;
- A simple scatterplot of price vs. Twitter mentions per day
- Potential problems:
 - Only 6 days of data
 - More granular financial data is not cheaply available
- We do not anticipate this model being very helpful
- But it is a good starting point!

What We Have

November_Rain



- Code to **unzip** a directory of .json files;
- A **dictionary** of words and their corresponding 'sentiment value' based on the McLougran data set;
 - About the data set. It has over 8,000 English words: positive, negative, or neutral in sentiment
- Code to return the **number of times** a specific word is mentioned and to find the stock ticker data using the yahoo_finance package;



- ✓ We now have all the working parts to create our base model!

Challenges We Have

November_Rain



➤ Scalability issues:

- 100GB data is a lot for our laptops
- The entire training data set
- We have compromised by reading 30 – 60 json files at a time (one hour of twitter data)
- ✓ **Possible solution:** to push the entire training set to AWS and run our training data there.



- **General programming issues:** (in particular type conversion issues)
 - We must determine the most efficient strategy with regards to data manipulation
 - Our approach: to deal exclusively with the data as either an RDD object or a pyspark dataframe, with lots of alternation between them
 - We have found: manipulation and transformation of these data types is not ideal
 - ✓ **Possible solution:** to perform any cleaning/filtering of the files while they are still in json format, since these can be read as strings or dictionaries, which are much more flexible data types.
 - ✓ We are not yet sure what effect this approach will have on scalability.



➤ Our general approach:

- 1) Filter out all irrelevant tweets and any superfluous information contained within them(non-english, does not mention one of our companies, etc), and store as a new .json (or bundle of jsons)
- 2) Pass a function that analyzes each tweet one at a time. This function will compare every word in the tweet to the sentiment dictionary we have created
- 3) Calculate a 'daily sentiment score' by tallying all scores from tweets mentioning a company. Plot closing price as a function of this daily sentiment score. Compare to base model.
- 4) If there is a correlation between price and sentiment score, use same approach on a test dataset to predict future closing price based on that day's 'daily sentiment score'



- If time permits we will fine tune our NLP approach. Although the loughran database is a great start, it is not extensive and categorizing words as either completely positive or negative is a crude approach to NLP.
- It would be great to apply some more advanced NLP techniques to our data.

Thank You!